

Parallel Sort Investigation – Lab Assignment

Student details:

Manghiuc Ioana, Informatica, 10LF223

Device configuration:

Laptop ASUS VivoBook S 15 M3502RA

Processor AMD Ryzen™ 9 6900HX

Base speed: 3.30 GHz

Cores: 8

Logical processors: 16

Integrated GPU: AMD Radeon™ 680M

RAM: 16GB

Summary:

2.1. Execution times, communication times, and processing times	2
2.1.1. Measure Execution Time.....	2
2.1.2. Measure Computation Time & Communication Overhead	3
2.1.3. Scalability Analysis.....	4
2.2. Discussion and analysis of results	8
2.2.1. Comparative Analysis of Sorting Methods	8
2.2.2. Discussion of Communication Overhead	9
2.2.3. Analysis of Computational Time and Bottlenecks	9
2.2.4. Speedup and Efficiency Evaluation	10
2.2.5. Possible Inefficiencies and Suggestions for Improvement	10

Evaluation and discussion of performance results

2.1. Execution times, communication times, and processing times

The execution and computation time, as well as the communication overhead, were measured as follows:

The **Bucket Sort**, **Odd-Even Sort**, and **Shell Sort** algorithms were tested using two, four, and eight (my device's maximum) cores. Then, for each core configuration the algorithm was tested three times to ensure data consistency. All these three algorithms were tested on an array of 1.000.000 elements and on an array of 10.000.000 elements, which are stored in two separate files.

The **Direct (Selection) Sort** was tested only on an array of 1.000.000 elements, using two, four, and eight cores. Each core configuration was tested three times.

The **Ranking Sort** algorithm was tested on an array of 1.000.000 elements using two, four, and eight cores. Each core configuration was tested once. I made this decision because the execution time with eight cores was already past the thirty-minute mark, and in my previous runs where I was only measuring the execution time, the result was similar, and therefore concluded the differences in minutes or seconds would do little to aid in the following analysis.

2.1.1. Measure Execution Time

To measure the total execution time of each parallel sorting algorithm, I used the `MPI_Barrier` and `MPI_Wtime` functions, in the following manner:

- **`MPI_Barrier(MPI_COMM_WORLD)`** was called before and after the sorting function to ensure that all processes are synchronized. This guarantees that the timing starts and ends consistently across all MPI ranks.
- **`MPI_Wtime()`** was used to record the wall-clock time in seconds. It provides a high-resolution timestamp for the current time, making it suitable for performance measurements in parallel applications.
- The **`computation_time`** and **`communication_time`** variables are updated inside the sorting function using additional calls to `MPI_Wtime()` to break down the execution time into computation vs. communication components. These values are passed by reference to the sorting function. The signature of the sorting function is identical for each algorithm.

The code is structured as:

```
MPI_Barrier(MPI_COMM_WORLD);
double computation_time = 0.0, communication_time = 0.0;
double execution_time = MPI_Wtime();

// Sorting function call
MPI_OddEvenSort(data, rank, size, computation_time, communication_time);
MPI_Barrier(MPI_COMM_WORLD);
execution_time = MPI_Wtime() - execution_time;
```

2.1.2. Measure Computation Time & Communication Overhead

To evaluate the performance of each parallel sorting algorithm, I measured both the **computation time** and the **communication overhead** separately using MPI_Wtime(). These times are calculated inside the sorting functions and accumulated in computation_time and communication_time variables passed by reference.

The computation time includes only the time spent on performing local calculations that typically include sorting local partitions (using std::sort or DirectSort), calculating ranks, maximum values, or bucket assignments, merging sorted chunks, local comparisons or evaluations in loops.

Each computational section is structured as follows:

```
start_time = MPI_Wtime();
// perform computation
computation_time += MPI_Wtime() - start_time;
```

Communication overhead was measured by wrapping all MPI communication calls with MPI_Wtime() and summing the durations. These include:

- MPI_Scatter / MPI_Gather / MPI_Gatherv: Distributing or collecting data
- MPI_Bcast: Broadcasting data to all processes (used in Ranking Sort)
- MPI_Sendrecv: Peer-to-peer communication (Odd-Even and Shell Sorts)
- MPI_Allreduce: Global status aggregation or finding max values
- MPI_Alltoallv / MPI_Alltoall: Complex many-to-many exchanges (Bucket Sort)

Each communication section is structured as follows:

```
start_time = MPI_Wtime();
MPI_<function>();
communication_time += MPI_Wtime() - start_time;
```

For all five sorting algorithms, the communication overhead has a significantly lower value in comparison to the computation time, which oftentimes was very close in value, if not identical, to the total execution time, meaning the time spent by MPI processes for communication does not affect the total time greatly. In the following scalability analysis, I will provide real data that I obtained in the measuring process for a better understanding.

2.1.3. Scalability Analysis

Sorting method	Data size	Cores	Execution nr.	Execution time	Computation Time	Communication time
Direct Sort	1.000.000	2	1.	1435.58	1435.34	0.238449
			2.	1311.69	1311.69	0.0012313
			3.	1342.55	1340.84	1.7028
		AVERAGE TIMES		1363.27	1362.62	0.64749
	1.000.000	4	1.	462.747	460.397	2.34866
			2.	362.982	360.679	2.31723
			3.	521.992	519.621	2.36876
		AVERAGE TIMES		449.240	446.899	2.344883
	1.000.000	8	1.	101.841	101.541	0.297848
			2.	124.074	122.728	1.34468
			3.	138.478	137.347	1.12929
		AVERAGE TIMES		121.4643	120.5386	0.9239

The execution time for Direct Sort increases as the number of MPI processes decreases, due to fewer cores sharing the computational workload. Although the use of MPI parallelism does improve performance to some extent, the algorithm's intrinsic time complexity of $O(n^2)$ severely limits its scalability. Here I must mention that from the task description (Direct Sort) it is not clear exactly which sorting algorithm it refers to, and it was clarified to me during a laboratory that it refers to the Selection Sort, which is why I implemented it, although I am still not sure it was required. Therefore, I acknowledge that using the `std::sort` algorithm instead would have improved Direct Sort's performance significantly.

Even though the initial partitioning and gathering of data are parallelized, the rest of the sorting work remains local to each process.

While testing on a smaller dataset of 1,000 elements, I observed that the collation (merge) operation was the most time-consuming. To optimize this part of the implementation, I used rvalue references for the input arrays, allowing them to be safely moved from. Additionally, I applied `std::move` to transfer ownership of elements rather than copying them, thereby reducing memory overhead and improving efficiency during the merge.

Despite this optimization, the algorithm still exhibits poor scalability as the number of processes decreases or the dataset size increases. This makes Direct Sort a weak candidate for scalable parallel execution when compared to other algorithms that distribute both computation and merging more evenly across processes

Bucket sort	1.000.000	2	1.	0.374329	0.369164	0.0025128
			2.	0.368792	0.364334	0.0023298
			3.	0.36999	0.364954	0.0023744
			AVERAGE TIMES	0.371037	0.36615	0.0024056
	10.000.000	2	1.	4.44484	4.40238	0.0249644
			2.	4.42754	4.38871	0.0195519
			3.	4.44286	4.40592	0.0178908
			AVERAGE TIMES	4.438413	4.399003	0.02080236
	1.000.000	4	1.	0.369947	0.360363	0.0072893
			2.	0.360507	0.354747	0.0034577
			3.	0.362048	0.35565	0.004141
			AVERAGE TIMES	0.3641673	0.35692	0.0049626
	10.000.000	4	1.	4.1921	4.15086	0.028771
			2.	4.189	4.15473	0.0214581
			3.	4.47079	4.43848	0.0195564
			AVERAGE TIMES	4.283963	4.248023	0.02326183
	1.000.000	8	1.	0.348963	0.342014	0.0049728
			2.	0.362234	0.355027	0.0050962
			3.	0.356924	0.344396	0.0105488
			AVERAGE TIMES	0.3560403	0.3471456	0.0068726
	10.000.000	8	1.	4.17184	4.11901	0.038731
			2.	4.13118	4.06617	0.0550711
			3.	4.13967	4.08679	0.0413536
			AVERAGE TIMES	4.147563	4.090656	0.0450519

The execution time for the Bucket Sort algorithm remains remarkably consistent across different numbers of MPI processes. For a dataset of 10 million elements, the execution time only varies slightly (e.g., between 4.10 and 4.45 seconds). This stability reflects efficient workload distribution through balanced value-based buckets.

The values are effectively partitioned into nearly balanced buckets, with each process having data within an allocated range (e.g., process 0 has data between 0 and 250 etc.), and local sorting is performed using `std::sort`, which is much faster than the Selection Sort used in Direct Sort. This combination ensures low computation time and strong scalability regardless of the number of processes.

Odd Even	1.000.000	2	1.	0.203022	0.169429	0.0284726
			2.	0.204297	0.170137	0.0271901
			3.	0.20222	0.168313	0.0288998
			AVERAGE TIMES	0.2031796	0.169293	0.0281875
	10.000.000	2	1.	2.29512	1.96192	0.282614
			2.	3.48884	2.98204	0.451242
			3.	2.37936	2.02593	0.278151
			AVERAGE TIMES	2.721106	2.323296	0.3373356
	1.000.000	4	1.	0.143749	0.0883962	0.0313318
			2.	0.136588	0.0816625	0.0245181
			3.	0.141281	0.0906397	0.0245264
			AVERAGE TIMES	0.1405393	0.08689946	0.0267921
	10.000.000	4	1.	1.61296	1.03227	0.289638
			2.	1.91337	1.12219	0.419559
			3.	1.61772	1.0686	0.275457
			AVERAGE TIMES	1.714683	1.074353	0.328218
	1.000.000	8	1.	0.126923	0.0533552	0.0265541
			2.	0.11382	0.0488368	0.0245451
			3.	0.115237	0.0470248	0.0296468
			AVERAGE TIMES	0.11866	0.04973893	0.0269153
	10.000.000	8	1.	1.21149	0.530941	0.27321
			2.	1.60533	0.823481	0.321144
			3.	1.20826	0.532383	0.291041
			AVERAGE TIMES	1.341693	0.628935	0.2951316

The compare and exchange Odd-Even algorithm has the best overall performance among the five implementations. Although the difference between execution times for different numbers of MPI processes is more noticeable, decreasing as the number of cores increases, it consistently outperforms the others in terms of speed.

The iterative compare-and-exchange pattern allows processes to refine their data ordering through minimal and structured communication, and the computation is optimized by using `std::merge` for the merging operation, and `std::copy`. As such, even with communication overhead it remains the fastest solution for both 1 million and 10 million datasets.

Ranking Sort	1.000.000	2	1.	7613.37	7613.36	0.0033229
		4	1.	3966.48	3966.47	0.003315
		8	1.	2337.14	2331.05	6.08485

The Ranking Sort method very noticeably delivers the worst performance out of all the five implementations. Due to its $O(n^2)$ time complexity, it scales poorly and remains inefficient even when parallelized. Although using eight MPI processes improves execution time compared to four or two, the total still reaches a decidedly inefficient 2337 seconds, which is abysmal even compared to the slowest Direct Sort run. This highlights the algorithm's fundamental inefficiency and poor suitability for large-scale parallel processing.

Shell Sort	1.000.000	2	1.	0.199478	0.193433	0.0016941
			2.	0.202363	0.189831	0.0017396
			3.	0.198994	0.19036	0.0016229
			AVERAGE TIMES	0.2002783	0.191208	0.00168553
	10.000.000	2	1.	2.32521	2.25006	0.015385
			2.	2.27525	2.20085	0.0146576
			3.	2.27381	2.20846	0.0157982
			AVERAGE TIMES	2.2914233	2.21979	0.01528026
	1.000.000	4	1.	0.133759	0.109682	0.0020478
			2.	0.138265	0.107101	0.002021
			3.	0.131723	0.108722	0.0022244
			AVERAGE TIMES	0.1345823	0.1085016	0.00209773
	10.000.000	4	1.	2.25899	1.94434	0.0245447
			2.	1.52009	1.28654	0.0233569
			3.	1.54834	1.28483	0.0240759
			AVERAGE TIMES	1.775806	1.505236	0.0239925
	1.000.000	8	1.	0.0885662	0.0612442	0.0025764
			2.	0.0922914	0.0617859	0.003225
			3.	0.0994363	0.0635539	0.0025386
			AVERAGE TIMES	0.0934313	0.0621946	0.00278
	10.000.000	8	1.	1.36737	1.00849	0.0325305
			2.	1.36252	1.09506	0.0313459
			3.	1.36553	1.11074	0.031535
			AVERAGE TIMES	1.36514	1.07143	0.0318038

Following very closely behind the Odd-Even implementation, the execution times recorded for the Shell Sort algorithm are the second most efficient. Its hybrid structure—using local `std::sort` followed by structured compare-and-merge operations—enables efficient parallel sorting with minimal communication overhead, notably lower than the communication overhead measured for the Odd-Even algorithm. While its execution time does increase as the number of MPI processes decreases, Shell Sort scales well overall and remains highly efficient even for larger datasets.

2.2. Discussion and analysis of results

2.2.1. Comparative Analysis of Sorting Methods

The implemented sorting algorithms show distinct differences in execution time and communication overhead. Ranked from best to worst in terms of performance:

- I. Odd-Even Sort - this algorithm consistently delivers the best execution time. Its parallelized compare-and-exchange pattern greatly reduces computation time, which is what makes this algorithm so efficient. However, this results in higher communication frequency when compared to the Shell or Bucket Sort. It scales well and remains fast across varying dataset sizes and numbers of MPI processes.
- II. Shell Sort – similarly, this algorithm leverages a parallelized compare-and-merge structure, supported by efficient operations such as `std::merge` and `std::copy`. Although second in overall performance, its communication overhead is noticeably lower than the communication overhead recorded for the Odd-Even algorithm. This indicates that the distinction lies mainly in Odd-Even’s faster computation rather than communication efficiency.
- III. Bucket Sort – while not as fast as Odd-Even or Shell Sort, this algorithm greatly outperforms the remaining two. Its design allows each process to sort values within predefined value ranges, leading to balanced workloads and efficient use of `std::sort` for local computation. It must also be pointed out that this algorithm had the most consistent performance across multiple runs using varying numbers of MPI processes, which reflects its stability. Its communication overhead is comparable to the Shell Sort overheads, further proving that the Odd-Even communication time is relatively high.
- IV. Direct Sort – the operation of repeatedly merging arrays diminishes this algorithm’s performance, which is worsened by the $O(n^2)$ time complexity of the selection sort algorithm, therefore rendering it inefficient for larger datasets. Its communication overhead is not comparable to the other sorting methods, proving to be less predictable, with values peaking at 2.344883 for four processes and dropping to 0.64749 for two processes, which highlights its instability under different configurations.
- V. Ranking Sort – decidedly the worst overall performance due to the inefficient ranking approach which results in $O(n^2)$ time complexity. Although the communication overhead increases significantly when eight cores are used, the algorithm’s fundamental issue lies in its high computational burden, thus making it unsuitable for large datasets.

2.2.2. Discussion of Communication Overhead

Communication time affects each algorithm differently based on its communication pattern and frequency of data exchange. Among all implementations, Odd-Even Sort is the most sensitive to communication latency. This is due to its iterative compareAndExchange mechanism, which involves frequent MPI_Sendrecv calls between neighboring processes in each sorting phase. While it achieves excellent overall speed, its communication time is noticeably higher than Shell Sort and Bucket Sort, especially on larger datasets. In contrast, Shell Sort and Bucket Sort show lower communication overheads. Shell Sort uses a similar neighbor-exchange pattern but fewer iterations due to its stride-based grouping. Bucket Sort benefits from structured MPI_Alltoallv and MPI_Gatherv communication, which, while complex, are performed in fewer steps. Direct Sort and Ranking Sort, while less affected by frequent communication, suffer from inefficiencies elsewhere, even though their communication times aren't particularly high. To optimize communication-heavy algorithms like Odd-Even, reducing the number of synchronization points could help lower communication latency.

2.2.3. Analysis of Computational Time and Bottlenecks

Each algorithm's performance is influenced heavily by its computational structure:

- Direct Sort is hindered by its use of Selection Sort, a $O(n^2)$ algorithm, which results in poor scalability. The merge step at the root process also becomes a bottleneck, especially as process count increases. However, replacing this computation technique with `std::sort` is what would reduce processing time significantly.
- Bucket Sort is computation-efficient due to its use of value-based partitioning and reliance on fast local sorting (`std::sort`). Its only bottleneck lies in computing and distributing buckets.
- Odd-Even Sort and Shell Sort use `std::sort` for initial local sorting, which is fast, but Odd-Even has additional iterative merge steps that increase total computation time. Still, the parallelized exchange-and-compare, and exchange-and-merge techniques greatly reduce the total computational time.
- Ranking Sort has the worst bottleneck: every process compares each element with all others during the ranking phase, leading to an enormous number of unavoidable comparisons and very high computation time, regardless of how data is split.

2.2.4. Speedup and Efficiency Evaluation

Theoretical speedup is calculated as:

$$S = \frac{T_s}{T_p}$$

where T_s is the execution time on a single processor and T_p is the time on a multiprocessor with p processes.

In practice, Odd-Even Sort shows near-linear speedup due to its well-parallelized structure, especially when moving from 2 to 8 cores. Shell Sort also shows good speedup, though diminishing slightly due to merge synchronization overheads.

Bucket Sort, while not scaling as dramatically, maintains stable execution times, showing more consistency than raw speedup.

Direct Sort and Ranking Sort fall far short of theoretical speedup. For Direct Sort, the sequential merge at the root nullifies parallel gains. For Ranking Sort, the entire algorithm is dominated by quadratic computation, which is not improved meaningfully by additional cores.

Therefore, deviations arise from algorithmic design, not MPI inefficiencies. Parallelizing both computation and merge steps achieves the best real-world speedup.

2.2.5. Possible Inefficiencies and Suggestions for Improvement

Several inefficiencies were identified across the implementations:

- Direct Sort: Uses inefficient local sorting and centralized merging. Although the merging operation is optimized for heavy data by using rvalue references and `std::move`, replacing Selection Sort with `std::sort` will undoubtedly improve its performance.
- Bucket Sort: Performs well but relies on a fixed bucket range distribution. Implementing adaptive bucket sizing based on data distribution could improve load balancing and prevent skewing.
- Odd-Even Sort: Though fast, it could benefit from non-blocking communication and asynchronous computation to reduce idle time during barriers.
- Ranking Sort: Suffers from $O(n^2)$ operations. It should be avoided in practical use and replaced with a more efficient parallel sort like Merge Sort or Bucket Sort.
- Shell Sort: Overall efficient, but its iterative merge logic may become redundant for well-distributed datasets. Adding an early termination check based on global sorted state could cut iterations.

In general, reducing synchronization points, using asynchronous communication, and leveraging more efficient local algorithms can significantly enhance parallel performance across all sorting methods.